

SELECTIVE ACTIVATION INTERVENTION FOR LOCALITY-AWARE MULTI-HOP FACTUAL ADAPTATION

Anonymous Authors

Affiliation withheld for the current draft

ABSTRACT

Factual adaptation has two objectives: make edited facts usable in multi-hop reasoning and avoid changing unrelated generations. Prompting and parameter editing can improve one objective while obscuring the other. We study Value-CQAG, a selective activation intervention that forms a contrast between an edited-fact prompt and the original question prompt, then gates that contrast with a content router. On a frozen 100-case, 300-question MQuAKE-CF confirmation pool, question-level generation rises from 2.00% to 38.67% on Qwen2.5-7B and from 4.00% to 51.33% on Llama-3.1-8B. Explicit facts remain a stronger accuracy control (95.0% and 93.0%), so our claim is selectivity rather than state-of-the-art accuracy. On a separate Qwen 10,000-pair audit, the routed system achieves 98.49% exact and 99.62% semantic locality with a 1.52% false-positive route rate. Failure analysis separates routing misses from old-answer and multi-hop entity errors. The result is a reversible, auditable adaptation path with a clear remaining bottleneck: multi-hop composition after a successful route.

Index Terms— large language models, factual adaptation, activation intervention, locality, multi-hop reasoning, generative AI

1. INTRODUCTION

Factual adaptation is not a single objective. A useful system must both propagate an edited fact through a multi-hop answer and abstain on questions that do not depend on that edit. These requirements are easy to conflate: a prompt containing the new fact can improve answerability, while a parameter or activation change can also alter unrelated generations. The distinction matters for rapid, reversible adaptation, where an edit should be removable without producing a globally changed model.

Existing approaches solve different parts of this problem. Parameter-editing methods locate and rewrite factual associations in feed-forward layers [1–4]; learned editors provide an alternative amortized formulation of the same update problem [5], while recent surveys organize the broader design space and its evaluation pitfalls [6]. External memories store an edit for later retrieval [7]; and retrieval or in-context methods supply explicit evidence at inference time [8, 9]. MQuAKE shows why single-hop success is insufficient: the edited fact must survive composition with other entities [10]. A single accuracy number therefore hides the important question of where the system failed: did it miss the edit, retain the old answer, or select the wrong entity after receiving the edit?

We study Value-CQAG as a controlled answer to that question. It forms an answer-independent contrast between an edited-fact prompt and the original question prompt in the model’s value-projection channel. A fixed content router decides whether to apply that contrast; rejected queries follow the original prompt path, while accepted queries receive a bounded prompt-only intervention before the same frozen decoder. This decomposition makes scope and effect measurable separately. Figure 1 summarizes the signal path.

Our contributions are:

- A parameter-free, gated activation operator that separates edit direction from edit scope and can be bypassed query by query.
- A controlled cross-model evaluation on Qwen2.5-7B and Llama-3.1-8B against explicit-fact, retrieval-plus-ICL, and MEMIT controls on the same frozen confirmation protocol.
- A 10,000-pair end-to-end locality audit and a failure decomposition showing that routing errors are distinct from downstream multi-hop entity errors.

2. METHOD

2.1. Problem setup

Let $\mathcal{E} = \{(s_i, r_i, o_i^{\text{new}})\}_{i=1}^K$ be a set of edited subject–relation–object facts and let q be a question whose answer may require several facts. The base model f_θ is frozen. We seek a conditional intervention $\Delta(q, \mathcal{E})$ such that

$$\hat{y} = f_\theta(q; \Delta(q, \mathcal{E})) \quad (1)$$

is consistent with the edited facts, while $f_\theta(q; 0)$ is retained for unrelated questions. Parameter editing instead produces one model $f_{\theta+\Delta\theta}$ for all subsequent inputs. Our target is therefore a conditional computation: change the signal only when the query is related to the edit.

2.2. Activation target and gated operator

Let $H_l^0(q) \in \mathbb{R}^{T \times d_l}$ denote the hooked self-attention value-projection activations at layer l when the original question prompt is encoded, and let $H_l^+(q, \mathcal{E})$ denote the corresponding activations for the same question preceded by a compact edited-fact instruction. We form the answer-independent contrast

$$D_l(q, \mathcal{E}) = H_l^+(q, \mathcal{E}) - H_l^0(q). \quad (2)$$

No target answer or generated completion enters D_l . The fact-conditioned pass is used to construct the contrast; generation itself follows the original question prompt, which permits a rejected query to bypass the edit completely. Let P_w

be a shape-preserving binary mask over the final w prompt positions, and let $g(q, \mathcal{E}) \in \{0, 1\}$ be the router decision. At selected layers $\mathcal{L}$, the operator is

$$\tilde{H}_l = H_l^0 + g(q, \mathcal{E}) \alpha P_w D_l, \quad l \in \mathcal{L}. \quad (3)$$

The confirmation configuration uses $\mathcal{L} = \{20, 22, 24, 26\}$, $w = 16$, and $\alpha = 5$, selected on development cases and frozen before confirmation. Equation 3 is applied only during prompt encoding; weights, caches after the prompt pass, and later autoregressive states are not edited. Direct explicit-fact prompting remains an accuracy upper control. The purpose of the contrast is different: it exposes an edit through a removable, query-gated activation path while decoding from the original question prompt.

The construction is related to recent activation steering and representation intervention methods [11–14], which show that useful task information can be exposed or shifted in hidden states without updating model weights. It differs in two ways: the direction is obtained from an explicit-fact contrast rather than a learned human-preference direction, and a content router gates the direction on a query-by-query basis.

2.3. Content-token routing

Each edit stores the source-question templates used to construct its directions. For a query q , the router lowercases text, extracts alphanumeric tokens, removes a fixed English stop-word list, and takes the maximum Jaccard overlap with a stored source question $s \in \mathcal{S}_{\mathcal{E}}$:

$$\rho(q, \mathcal{E}) = \max_{s \in \mathcal{S}_{\mathcal{E}}} \frac{|T(q) \cap T(s)|}{|T(q) \cup T(s)|}. \quad (4)$$

The gate is

$$g(q, \mathcal{E}) = \mathbb{I}[\rho(q, \mathcal{E}) \geq \tau], \quad \tau = 0.35. \quad (5)$$

The threshold is fixed on development data and is not tuned on the confirmation pool. A router miss is counted separately from a generation error after a successful route. This abstention mechanism motivates reporting coverage and generation separately rather than hiding both behind one accuracy score.

2.4. First-order effect and locality identity

For a fixed answer position, let $m(q) = z_{y_{\text{new}}}(q) - z_{y_{\text{old}}}(q)$ be the new-versus-old logit margin. A local expansion around the unmodified activations gives

$$\Delta m \approx \alpha g \sum_{l \in \mathcal{L}} \langle P_w \nabla_{H_l} m, P_w D_l \rangle + O\left(\alpha^2 \sum_l \|P_w D_l\|_F^2\right). \quad (6)$$

The equation is diagnostic, not a global guarantee: a useful contrast must align with the local answer-margin gradient, whereas an unrelated equal-norm direction has no systematic alignment under an isotropic null model. Multi-token generation can still fail after a favorable local change, which motivates the separate old-answer and entity-error audit.

The router also gives an exact locality decomposition. For an unrelated query u , let $\phi = \Pr[g(u, \mathcal{E}) = 1]$ be the false-positive route rate and let $\eta = \Pr[\text{output changes} \mid g = 1]$. Because rejected queries bypass Equation 3,

$$L_{\text{exact}} = 1 - \phi\eta. \quad (7)$$

On the 10,000-pair audit, $\phi = 152/10,000 = 1.52\%$ and $\eta = 0.9934$, yielding $L_{\text{exact}} = 98.49\%$. Equation 7 explains why a high router rejection rate is necessary but not sufficient: a false route is highly likely to alter the output.

2.5. Two-stage prompt-pass implementation

Figure 1 shows the two stages. Stage 1 forms H_l^0 and H_l^+ with frozen passes and evaluates the router. Stage 2 either forwards the original activations or applies their contrast before the shared frozen decoder. The pictorial input symbols are schematic; all reported inputs are text. Extending the intervention through the first two generated tokens reduces question accuracy by 4.00 points on Qwen and 7.00 points on Llama, supporting the prompt-only design.

3. RELATED WORK

Knowledge editing has been studied through localized parameter updates such as ROME and MEMIT [2, 3], and through external edit memories such as SERAC [7]. The MQuAKE benchmark specifically exposes failures that only appear when several edited facts must be composed [10]. Our intervention leaves the parameters and external memory unchanged and instead controls the activation signal presented to a frozen decoder.

Activation steering methods add directions at inference time, including unsupervised activation addition [11], inference-time interventions for truthful behavior [12], and function-vector analyses of in-context computation [13]. Work on latent knowledge probing further suggests that internal representations can carry answer-relevant information even when the final generation is unreliable [14]. These methods motivate treating hidden states as a controllable signal, but they generally assume a direction is applied whenever a prompt is evaluated. Our contribution is the combination of an explicit-fact contrast, a lexical gate, and a separate end-to-end locality audit. Retrieval-augmented generation remains a strong accuracy control [8]; we therefore report it rather than presenting activation intervention as an accuracy replacement.

4. EXPERIMENTAL PROTOCOL

4.1. Models and data

We evaluate Qwen2.5-7B-Instruct and Llama-3.1-8B-Instruct. All principal comparisons use the same frozen 100-case MQuAKE-CF confirmation pool, containing 300 questions (three per case); MQuAKE is designed to test composition rather than isolated single facts [10]. The pool is separate from a reserved set that is not used in this study. For parameter editing, five deterministic unrelated questions per edited case form a 500-output locality check; the routed audit is reported separately below.

4.2. Baselines and metrics

We compare four adaptation modes: (i) the unmodified base model, (ii) explicit facts placed in the prompt, (iii) a transparent retrieval-plus-ICL control inspired by IKE [9], and (iv)




Fig. 1. Two-stage Value-CQAG pipeline. Stage 1 encodes question-only and edited-fact-conditioned prompts to obtain $H_l^0(q)$ and $H_l^+(q, \mathcal{E})$, while the router tests edit relevance. In Stage 2, rejected queries retain $H_l^0(q)$; accepted queries receive the contrast $D_l(q, \mathcal{E})$ through $\tilde{H}_l$ before frozen decoding.

EasyEdit MEMIT [3]. The retrieval corpus contains 500 non-overlapping development cases and uses MiniLM semantic retrieval with eight demonstrations. This is an IKE-style control, not a claim of an unmodified official IKE reproduction.

We report question-level generation accuracy, case-level generation accuracy (all questions in a case correct), new-answer preference, and exact-output locality. A generation hit requires normalized exact match to the new answer or a recorded alias. The routed locality audit crosses 100 edited cases with 100 unrelated questions, yielding 10,000 pairs. Exact locality requires an unchanged output string; semantic locality requires preservation of the task’s answer-equivalence status. The MEMIT locality protocol uses 500 fixed unrelated outputs and is not pooled with the routed audit. Confidence intervals are paired case-bootstrap intervals where reported.

4.3. Reproducibility

For Value-CQAG, models use 4-bit NF4 loading with fp16 compute, greedy decoding, and at most 24 new tokens; a generation hit is normalized exact match to the new answer or a recorded alias. For MEMIT, second-moment statistics use the local ‘wikimedia/wikipedia’ 20231101.en mirror, 100,000 samples, float32, and layers 4–8 of ‘mlp.down_proj’. Due to the 24-GB RTX 3090 memory limit, MEMIT models are loaded in fp16; the covariance linear system is solved in CPU double precision and only the resulting update is moved to the GPU. The scripts, hparams, frozen IDs, and raw JSON artifacts accompany this paper draft.

Table 1. Question-level and case-level generation accuracy on the same frozen 100-case (300-question) confirmation pool. Retrieval+ICL[†] was run on Qwen only, using eight demonstrations from 500 development cases.

Model	Method	Question	Case
Qwen2.5-7B	Base	2.00%	0.0%
	Value-CQAG	38.67%	18.0%
	Explicit facts	95.00%	90.0%
	Retrieval+ICL [†]	95.00%	93.0%
Llama-3.1-8B	Base	4.00%	0.0%
	Value-CQAG	51.33%	32.0%
	Explicit facts	93.00%	87.0%

5. RESULTS

5.1. Cross-model comparison

Table 1 states the paper’s central trade-off. Value-CQAG substantially improves over the base model on both architectures, but explicit facts and retrieval remain stronger on answer accuracy. The contribution is therefore controlled adaptation: the method changes the original prompt path only when the router accepts the query, rather than claiming to replace a long evidence prompt on raw accuracy.

Relative to the base model, the paired question-level gain is +36.67 percentage points for Qwen (95% case-bootstrap CI [29.33, 44.33]) and +47.33 points for Llama (95% CI [39.33, 55.33]). Retrieval-plus-ICL reaches 95.0% question-level and 93.0% case-level accuracy on Qwen, but its exact and semantic locality are only 25.0% and 61.4%, respectively, under its

Table 2. EasyEdit MEMIT on the frozen 100-case pool. Locality uses 500 unrelated outputs.

Model	Pref.	Q-gen.	Case	Exact loc.	Edit (s)
Qwen2.5-7B	16.0%	16.0%	8.0%	85.2%	52.41
Llama-3.1-8B	18.0%	14.33%	5.0%	73.4%	29.64

Table 3. Completed auxiliary evaluations. “Locality” is mismatched-direction locality for Value-CQAG and exact-output locality for ROME. Results are not pooled across protocols.

Setting	n	Q-gen.	Case	Locality	Seeds
Qwen CF, Value-CQAG	30×3	27.78 ± 2.22	12.22 ± 1.92	80.00 ± 14.53	3
Qwen T, Value-CQAG	16×3	21.53 ± 2.41	0.00	64.58 ± 7.22	3
Qwen2.5-3B CF	30	3.33	0.00	—	1
Llama CF, question-specific	80	42.92	28.75	—	1
Qwen CF, ROME	80	19.17	13.75	80.75	1

separate long-context protocol. These results position Value-CQAG as a selective, lower-interference alternative, not as the strongest answer-generation baseline.

5.2. Locality and parameter editing

The end-to-end routed audit on Qwen gives a 1.52% false-positive route rate, 98.49% exact-output locality, and 99.62% semantic locality. Conditional on a false route, exact locality falls to 0.66%, showing that the router is the main safety boundary rather than evidence that the injected direction is intrinsically local. The audit records 37 newly corrupted originally-correct outputs and only one newly fixed originally-incorrect output.

Table 2 compares the two MEMIT runs. The method has lower generation accuracy than Value-CQAG and substantially lower locality on Llama. These results indicate that a weight edit can provide a persistent change, but persistence is not equivalent to selective behavior on unrelated questions.

5.3. Failure decomposition

Among 300 Qwen questions, 116 are correct, 133 are other-entity or reasoning errors, 31 retain the old answer, 19 are router misses, and one is a refusal. Llama yields 154 correct, 96 other-entity or reasoning errors, 30 old-answer residuals, 19 router misses, and one refusal. The identical router-miss count across models separates coverage from the downstream generation bottleneck: after a successful route, multi-hop entity selection remains the dominant failure mode.

5.4. Robustness and auxiliary evaluations

Table 3 reports auxiliary evaluations. Their case pools and filtering rules differ from the confirmation protocol, so they are used only to test whether the observed behavior is consistent across settings, not as pooled evidence.

The explicit-fact control reaches 95.0% / 90.0% question/case accuracy on Qwen and 93.0% / 87.0% on Llama. On the earlier 30-case split it reaches 78.89% on Qwen2.5-7B and 74.44% on Qwen2.5-3B, confirming that the evaluation is

solvable and that the remaining gap is methodological rather than an impossible target.

5.5. Ablations and efficiency

The development and probe runs provide directional evidence for the operator in Equation 3. A single token gave 25% new-answer preference on a 12-question probe, whereas a 4–16-token window reached 75%. On the 30-case CF split, one layer (22) gave 46.67% preference and 76.67% mismatched locality; three layers (22,24,26) gave 66.67% preference but 63.33% locality. Continuing the injection for two decoding steps reduced confirmation question accuracy from 38.67% to 34.67%; a full token trajectory reached only 17.0% on a development validation. These are ablations, not additional independent confirmation claims.

For the Qwen 80-case efficiency audit, vector construction averaged 2.273 s/case, base generation 0.388 s/pair, and steered generation 0.924 s/pair. No-router exact locality was 3.33%; the router costs 3.584 s/case in the 10,000-pair audit.

6. DISCUSSION AND LIMITATIONS

The evidence supports a focused conclusion. Conditional activation intervention improves multi-hop factual adaptation on two instruction-tuned models without changing their parameters, while the router yields high system-level locality. The same audit also shows the boundary of that claim: after a false route, exact locality falls to 0.66%, and after a successful route, multi-hop entity selection remains the dominant failure mode. Retrieval is a strong accuracy control; Value-CQAG is useful when selective, reversible behavior is more important than maximizing answer accuracy.

The confirmation pool contains 100 cases and is not a broad pretraining-scale benchmark. The router uses lexical overlap with stored source-question templates and can miss paraphrases. The five-question parameter-editing locality protocol and 10,000-pair routed audit measure related but different quantities. Finally, the MEMIT comparison uses a documented hardware compatibility path (fp16 loading and CPU double solve) on a 24-GB GPU; it is not presented as an unmodified official MEMIT run. Future work should test relation-macro accuracy, independent seeds, semantic routing, matched end-to-end resource budgets, and a learned locality-regularized router. These limitations bound the claim; they do not change the observed separation between routing coverage and downstream composition.

7. CONCLUSION

We presented Value-CQAG, a selective activation intervention for multi-hop factual adaptation. Its central design choice is to separate what changes from where it changes: an explicit-fact contrast supplies the direction, while a content router controls exposure and rejected queries retain the original path. Across Qwen and Llama, this improves over the base model but does not surpass explicit prompting on raw accuracy; the value of the method is the auditable locality boundary and the resulting failure decomposition. The next technical target is clear: replace lexical routing and improve multi-hop composition without sacrificing that boundary.

Acknowledgment

This work received no external funding. The authors thank the open-source model and benchmark communities for releasing the software and data used in this study.

Compliance with Ethical Standards

This study uses publicly available benchmark cases and pre-trained models; it involves no human subjects, user data, or private or personally identifiable information. The evaluation is automated and reports aggregate outcomes. The reserved 60-case pool remains sealed and was not read or evaluated during this study.

8. REFERENCES

- [1] Mor Geva, Roei Schuster, Jonathan Berant, and Omer Levy, “Transformer feed-forward layers are key-value memories,” in Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing, 2021, pp. 5484–5495.
- [2] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov, “Locating and editing factual associations in GPT,” in Advances in Neural Information Processing Systems, 2022.
- [3] Kevin Meng, Arnab Sen Sharma, Alex Andonian, Yonatan Belinkov, and David Bau, “Mass-editing memory in a transformer,” in International Conference on Learning Representations, 2023.
- [4] Damai Dai, Li Dong, Yaru Hao, and Zhifang Sui, “Knowledge neurons in pretrained transformers,” in Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), 2022, pp. 8493–8502.
- [5] Nicola De Cao, Wilker Aziz, and Ivan Titov, “KnowledgeEditor: Editing knowledge in language models,” in International Conference on Learning Representations, 2021.
- [6] Yunzhi Yao et al., “Editing large language models: Problems, methods, and opportunities,” arXiv preprint arXiv:2305.13172, 2023.
- [7] Eric Mitchell, Charles Lin, Antoine Bosselut, Chelsea Finn, and Christopher D. Manning, “Memory-based model editing at scale,” in Proceedings of the 39th International Conference on Machine Learning, 2022, pp. 15817–15831.
- [8] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Kuttler, Mike Lewis, Wen tau Yih, Tim Rocktaeschel, Sebastian Riedel, and Douwe Kiela, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in Advances in Neural Information Processing Systems, 2020.
- [9] Ce Zheng, Lei Li, Qingxiu Dong, Yuxuan Fan, Zhiyong Wu, Jingjing Xu, and Baobao Chang, “Can we edit factual knowledge by in-context learning?,” in Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, 2023, pp. 4862–4876.
- [10] Zexuan Zhong, Yifan Wu, Michael D. Morse, Yiran Zhang, Steven H. Wang, et al., “MQuAKE: Assessing knowledge editing in language models via multi-hop questions,” in Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing, 2023.
- [11] Alexander Turner, Lisa Thiergart, David Udell, Gavin Leech, Ulisse Mini, and Mark Goldstein, “Activation addition: Steering language models without optimization,” arXiv preprint arXiv:2308.10248, 2023.
- [12] Kang Li, Yutong Zhang, Ziqian Zhong, and Liwei Wang, “Inference-time intervention: Eliciting truthful answers from a language model,” in Advances in Neural Information Processing Systems, 2023.
- [13] Eric Todd, Millicent L. Li, Arnab Sen Sharma, Aaron Mueller, Byron C. Wallace, and David Bau, “Function vectors in large language models,” in International Conference on Learning Representations, 2024.
- [14] Collin Burns, Haotian Ye, Dan Klein, and Jacob Steinhardt, “Discovering latent knowledge in language models without supervision,” in International Conference on Learning Representations, 2024.